

Netflix Movie Recommendation Engine



Netflix held the **Netflix Prize open competition** for the **best algorithm** to **predict user ratings** for **films**. The **grand prize** was **\$1,000,000**. This is the **dataset** that was used in that competition and we're going to use it to build an effective **recommendation engine**.

There's a **movie_titles.csv** file containing the **details** of the **movie** and there are 4 other **combined_data_(1,2,3,4).txt** files containing the **user ratings**.

Loading the dependencies and the data

```
In [1]: # importing basic dependencies
import numpy as np
import pandas as pd
import scipy as sp
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
import os
warnings.filterwarnings('ignore')
```

```
In [2]: # loading the movie_titles.csv
movies=pd.read_csv('movie_titles.csv', names=['id','year','title'])

# peek into the dataframe
movies.head()
```

Out[2]:

	id	year	title
0	1	2003.0	Dinosaur Planet
1	2	2004.0	Isle of Man TT 2004 Review
2	3	1997.0	Character
3	4	1994.0	Paula Abdul's Get Up & Dance
4	5	2004.0	The Rise and Fall of ECW

```
In [3]: movies.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17770 entries, 0 to 17769
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   id      17770 non-null     int64
1   year    17763 non-null     float64
2   title   17770 non-null     object
dtypes: float64(1), int64(1), object(1)
memory usage: 416.6+ KB
```

In the following piece of code we're **combining** all the **four text files** containing the **user ratings** into one massive **rating.csv file**.

```
In [4]: # creating a new file if it doesn't exist already
if not os.path.isfile('rating.csv'):

    # opening the newly created file on the 'write' mode
    rating=open('rating.csv', mode='w')

    separated_data=['combined_data_1.txt','combined_data_2.txt',
                    'combined_data_3.txt','combined_data_4.txt']

    # iterating through the separate files containing the user ratings
    for file in separated_data:
        with open(file) as f:
            # processing each line of the currently open file
            for line in f:
                line=line.strip()
                if line[-1]==':':
                    movie_id=line[:-1]
                else:
                    combining_rows=[x for x in line.split(',')
                                     combining_rows=[movie_id]+combining_rows
                                     rating.write(','.join(combining_rows))
                                     rating.write('\n')

            f.close()
    rating.close()
```

```
In [5]: # loading the above new csv file into a dataframe
rating=pd.read_csv('rating.csv', sep=',',
                  names=['Movie_Id', 'Customer_Id', 'Rating', 'Date'])
```

```
In [6]: # peek into the dataframe
rating.head()
```

Out[6]:

	Movie_Id	Customer_Id	Rating	Date
0	1	1488844	3	2005-09-06
1	1	822109	5	2005-05-13
2	1	885013	4	2005-10-19
3	1	30878	4	2005-12-26
4	1	823519	3	2004-05-03

```
In [7]: # merging the movies dataframe with rating dataframe
df=pd.merge(movies, rating, left_on='id', right_on='Movie_Id', how='inner')
```

```
In [8]: # peek into the dataframe
df.head()
```

Out[8]:

	id	year	title	Movie_Id	Customer_Id	Rating	Date
0	1	2003.0	Dinosaur Planet	1	1488844	3	2005-09-06
1	1	2003.0	Dinosaur Planet	1	822109	5	2005-05-13
2	1	2003.0	Dinosaur Planet	1	885013	4	2005-10-19
3	1	2003.0	Dinosaur Planet	1	30878	4	2005-12-26
4	1	2003.0	Dinosaur Planet	1	823519	3	2004-05-03

Clearly we can drop **id** as it is merely as copy of **Movie_Id** and the **title**, **year** and **Date** are also not relevant in the context of building a **recommendation system**.

We're going to work with the following three variables only **Movie_Id**, **Customer_Id** and **Rating**.

```
In [9]: # dropping the unnecessary variables
df=df[['Movie_Id', 'Customer_Id', 'Rating']]
```

Exploratory Data Analysis

```
In [10]: # checking the metadata
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 100480507 entries, 0 to 100480506
Data columns (total 3 columns):
#   Column      Dtype
---  -
0   Movie_Id    int64
1   Customer_Id int64
2   Rating      int64
dtypes: int64(3)
memory usage: 3.0 GB
```

```
In [11]: print(f'There are {df.Movie_Id.unique().shape[0]} unique movies in the dataframe')
```

There are 17770 unique movies in the dataframe

```
In [12]: print(f'There are {df.Customer_Id.unique().shape[0]} customers in the dataframe')
```

There are 480189 customers in the dataframe

```
In [13]: print(f"The minimum Customer_Id and the maximum Customer_Id in the df are {df.Customer_Id.min()} and {df.Customer_Id.max()}")
```

The minimum Customer_Id and the maximum Customer_Id in the df are 6 and 2649429 respectively and we know that there are only 480189 unique customers in the df which will lead to confusion. Hence we're going to map them as continuous integers

```
In [14]: # creating a dict map to map old Customer_Id to new and continuous Customer_Id
mapping={old:new for old,new in
         list(zip(sorted(df.Customer_Id.unique()),
                  range(df.Customer_Id.unique().shape[0])))}
```

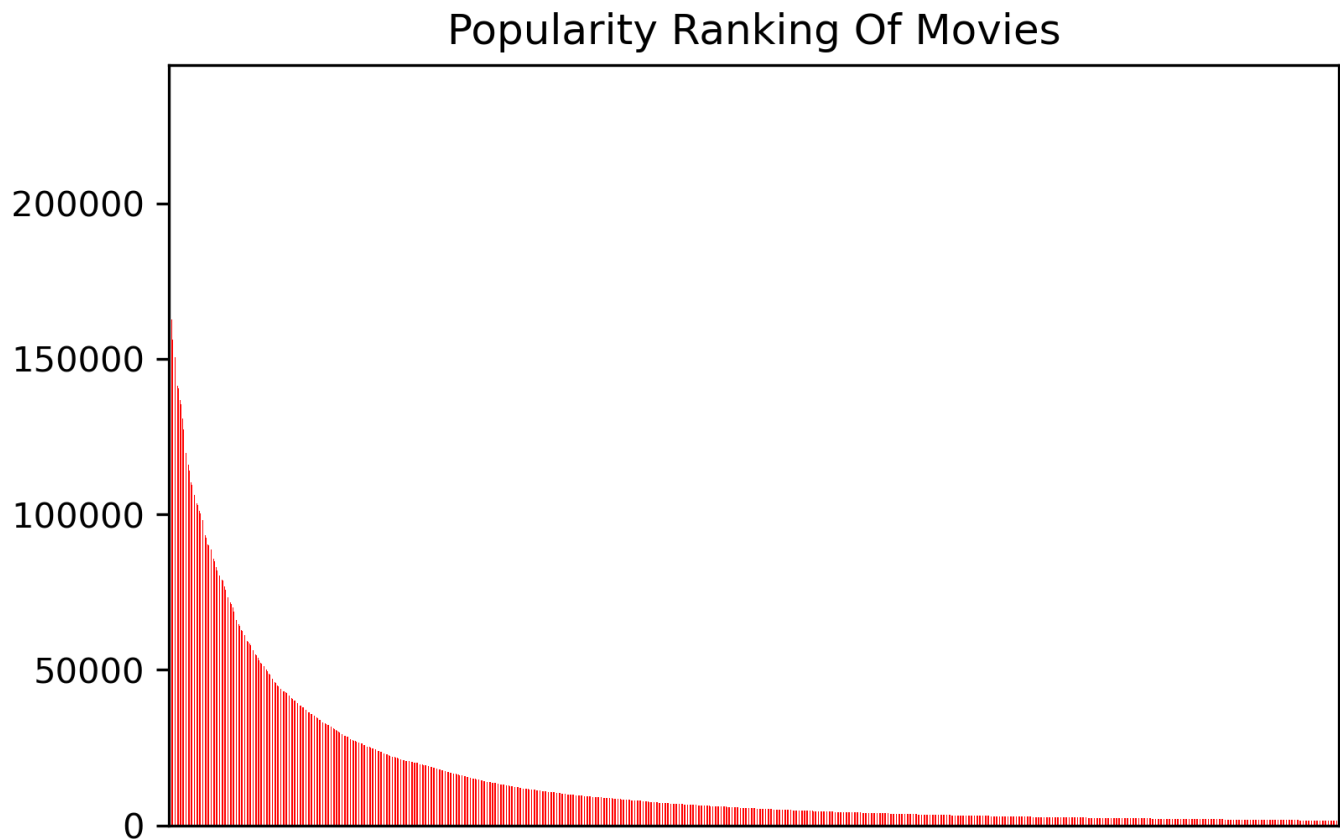
```
In [15]: # applying the above mapping
df['Customer_Id']=df['Customer_Id'].map(mapping)
```

```
In [16]: # checking if the new Customer_Id is continuous
print(df.Customer_Id.unique().shape[0])
print(df.Customer_Id.min())
print(df.Customer_Id.max())
```

```
480189
0
480188
```

Checking the **distribution** of **popularity** of the **movies**


```
In [17]: plt.figure(dpi=300)
plt.title('Popularity Ranking Of Movies')
df.Movie_Id.value_counts()[:6000].plot(kind='bar', color='red')
plt.xticks(ticks=[]);
```



The above plot shows the **long tail distribution** of the movies. The first **few 100 movies** or so have ratings **steeply dropping** from **230000** to **50000** and the rest of the movies have very low number of ratings. (Note that we've considered only the 6000 most popular movies).

Checking the **distribution** of **ratings** of the **movies**

```
In [18]: plt.figure(dpi=300)
plt.title('Distribution of Ratings')
sns.countplot(df.Rating, palette='Reds');
```



As we can clearly see that **4 stars** and **3 stars** are the **most** and the **second most frequent ratings** given by users.

```
In [19]: # there are no NaNs in any of the columns
df.isna().sum()
```

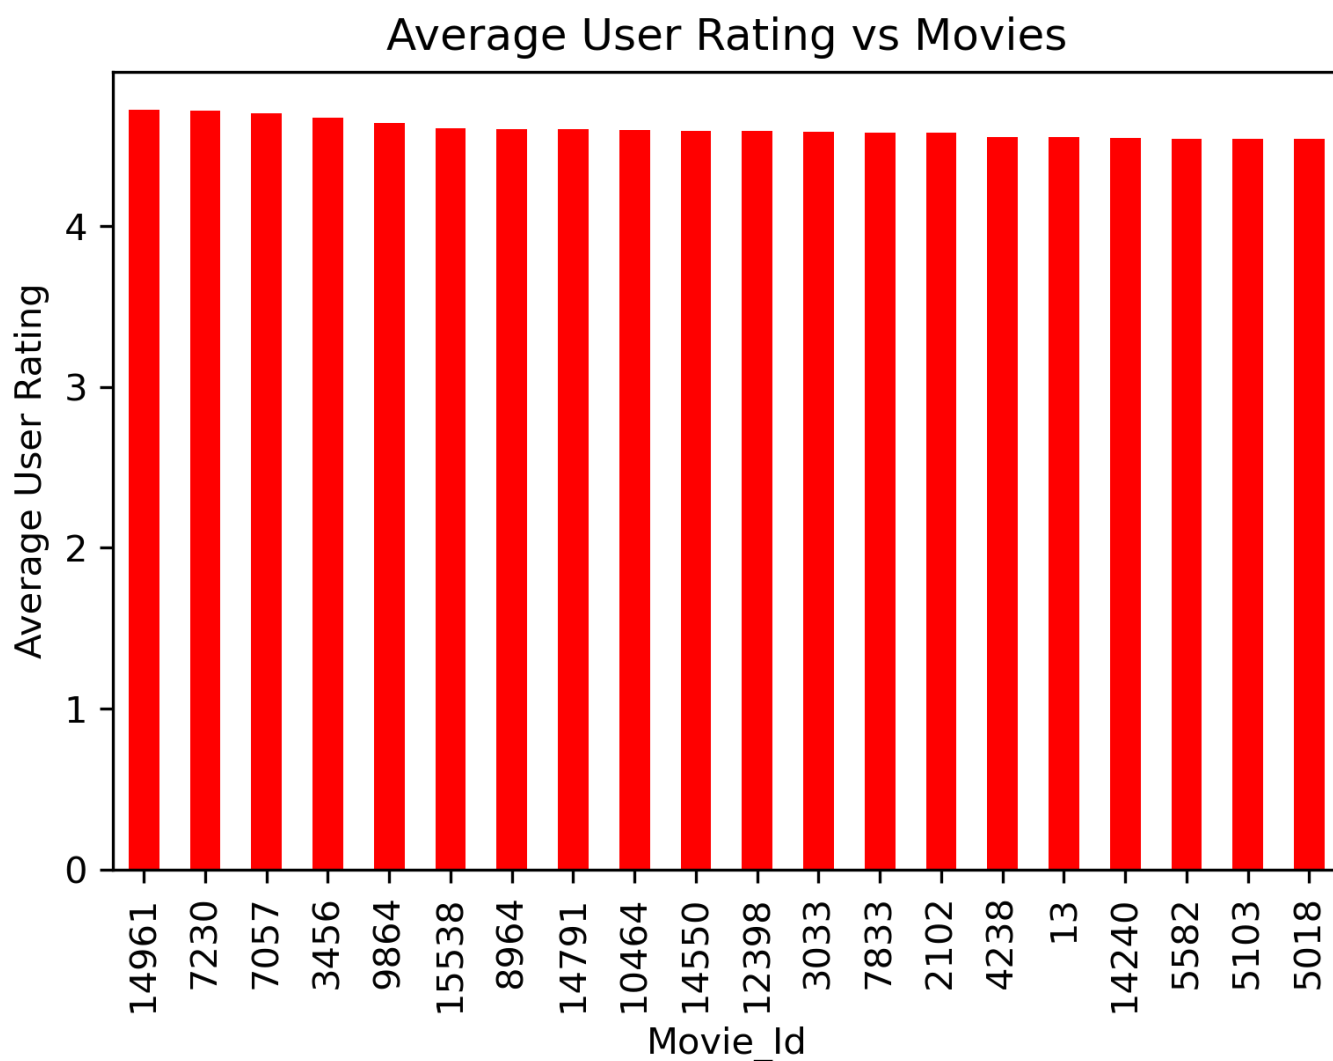
```
Out[19]: Movie_Id      0
Customer_Id    0
Rating         0
dtype: int64
```

```
In [20]: # there are no duplicated ratings
df.duplicated(['Movie_Id', 'Customer_Id']).sum()
```

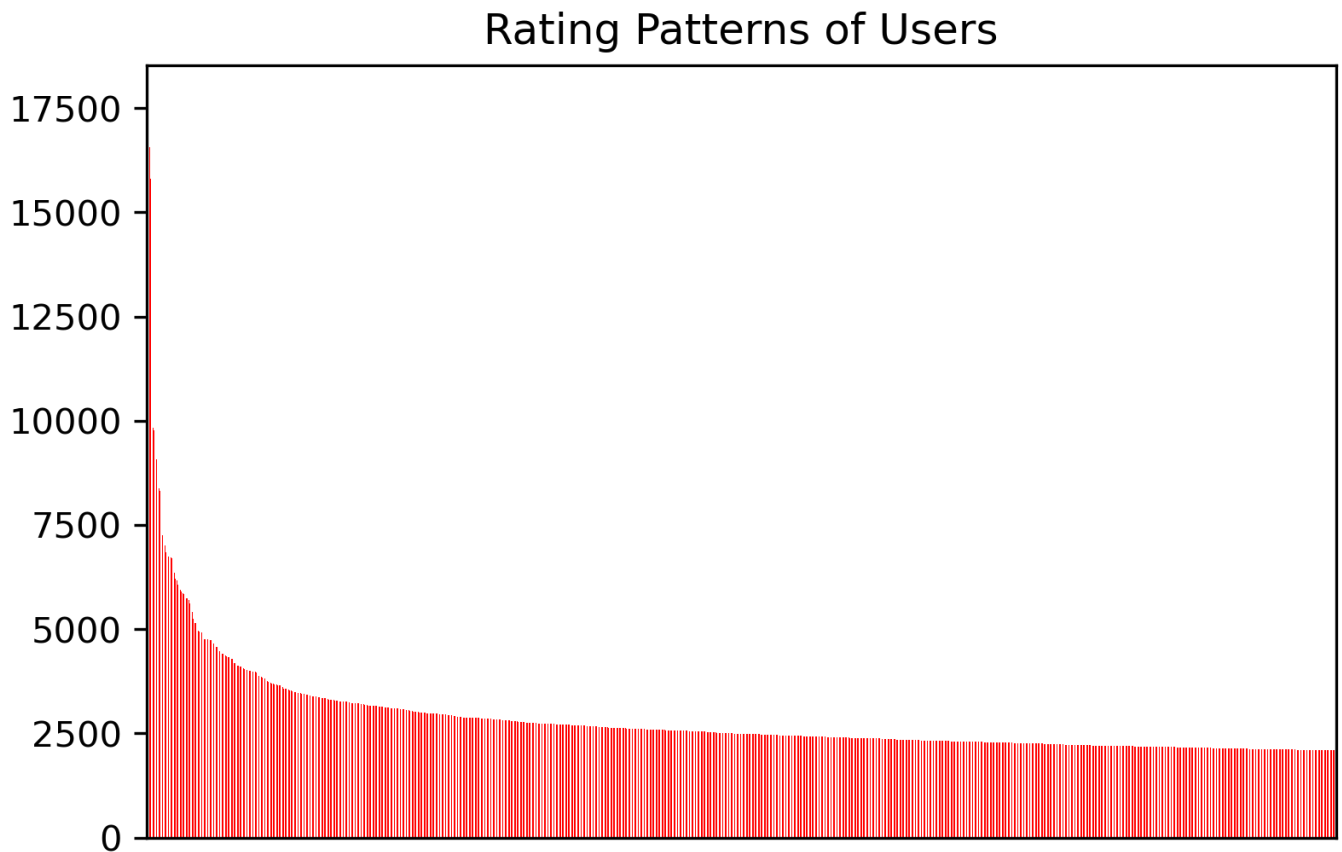
```
Out[20]: 0
```

List of **top 20 best** movies as per the **average user rating**.

```
In [21]: # average user rating of top 20 movies
plt.figure(dpi=300)
plt.title('Average User Rating vs Movies')
plt.ylabel('Average User Rating')
(df.groupby(by='Movie_Id')['Rating']
 .mean().sort_values(ascending=False)[:20]).plot(kind='bar', color='red');
```



```
In [22]: plt.figure(dpi=300)
plt.title('Rating Patterns of Users')
df.Customer_Id.value_counts()[1000].plot(kind='bar', color='red')
plt.xticks(ticks=[]);
```



The above plot displays the **rating patterns** of the **top 1000 frequent users** and their rating behavior ranges from rating almost all movies to rating only handful of movies

```
In [23]: # statistical description of ratings
df.Rating.describe()
```

```
Out[23]: count      1.004805e+08
mean        3.604290e+00
std         1.085219e+00
min         1.000000e+00
25%         3.000000e+00
50%         4.000000e+00
75%         4.000000e+00
max         5.000000e+00
Name: Rating, dtype: float64
```

Utility Matrix

Here we're going to produce the **utility matrix** by **pivoting** the **dataframe df** with **Movie_Id** as **rows**, **Customer_Id** as **columns** and the **Rating** as the cell **values**. Also we're going to use **item-item cosine similarity based collaborative filtering** technique to build our **recommendation system**.

Let's **centre** the **rating vectors** of every **movie** by **subtracting** it by **average rating** of the movie. We're doing this in order to make better sense of the **cosine similarity**.

```
In [24]: # finding the average rating for each movie
mean_rating=df.groupby(by='Movie_Id')[['Rating']].mean()
mean_rating.head()
```

Out[24]:

	Rating
Movie_Id	
1	3.749543
2	3.558621
3	3.641153
4	2.739437
5	3.919298

```
In [25]: # merging the above df with the original df
tmp=pd.merge(df, mean_rating, left_on='Movie_Id',
              right_index=True, how='inner')
tmp.head()
```

Out[25]:

	Movie_Id	Customer_Id	Rating_x	Rating_y
0	1	270045	3	3.749543
1	1	149546	5	3.749543
2	1	160878	4	3.749543
3	1	5466	4	3.749543
4	1	149791	3	3.749543

```
In [26]: # creating a df containing the movie-wise centred ratings
df=tmp[['Movie_Id','Customer_Id']]
df['Rating']=tmp['Rating_x']-tmp['Rating_y']
df.head()
```

Out[26]:

	Movie_Id	Customer_Id	Rating
0	1	270045	-0.749543
1	1	149546	1.250457
2	1	160878	0.250457
3	1	5466	0.250457
4	1	149791	-0.749543

Since **utility matrix** is going to be a massive, we're going to use **csr_matrix** to represent it.

```
In [27]: rows=df.Movie_Id.values  
cols=df.Customer_Id.values  
data=df.Rating.values  
utility=sp.sparse.csr_matrix((data, (rows, cols)),  
shape=(1+df.Movie_Id.unique().shape[0],1+df.Customer_Id.unique().shape[0]))  
utility
```

```
Out[27]: <17771x480190 sparse matrix of type '<class 'numpy.float64'>'  
with 100480507 stored elements in Compressed Sparse Row format>
```

```
In [28]: # unsupervised learning algo for recommendation system  
from sklearn.neighbors import NearestNeighbors
```

```
In [29]: # capturing the pair of movie index and user index who actually rated the movie  
mov_idx,user_idx=utility.nonzero()
```

Recommender System Based On Collaborative Filtering Item-Item Cosine Similarity

```

In [30]: def recommendations(user_id):
# list of top 5 movies rated by the given user
user_top_movies=utility[:,user_id].toarray().reshape(-1,)

# finding the argmax for top 5 ratings
user_top_movies=user_top_movies.argsort()[-5:][::-1]

print(f'The list of top 5 rated movies by the Customer_Id {user_id} are')
for film in movies.iloc[user_top_movies]['title'].values:
    print(film)

# rating vectors of all the movies rated by this user
# in otherwords utility submatrix
sub_utility=utility[mov_idx[user_idx==user_id]]

# instantiating KNN model with default 5 neighbors
knn=NearestNeighbors(metric='cosine', n_jobs=-1, algorithm='brute')

# training the above algorithm on the utility submatrix
knn.fit(sub_utility)

# collecting list of movies not watched by this user
unwatched
=list(set(range(1,17771)).difference(set(mov_idx[user_idx==user_id])))

unwatched_rating=[-10]*17771
# for every unwatched film by this user
for newfilm in unwatched:

    # we're finding the rating vector of the movie
    rate_vector=utility[newfilm].toarray()

    # extracting the cosine distances and movie_id
    distance,movieid=knn.kneighbors(rate_vector, 5, return_distance=True)

    # reshape the distance and the ids
    distance,movieid=distance.reshape(-1,),movieid.reshape(-1,)

    # converting cosine distance to cosine similarity
    similarity=1-distance

    # ratings of the aforementioned 5 most similar movies by this user
    rates=utility[:,user_id][movieid].toarray().reshape(-1,)

    # finding the weighted mean of the 5 highly rated movies
    predicted_rating=np.dot(similarity,rates)/np.abs(similarity).sum()

    # assigning the predicted_rating at the respective index
    # corresponding to the movie index
    unwatched_rating[newfilm]=predicted_rating

# converting it to numpy array
unwatched_rating=np.array(unwatched_rating)

# finding the argmax for top 5 ratings
top_recommends=unwatched_rating.argsort()[-5:][::-1]

print(f'\nThe list of top 5 recommended movies for the Customer_Id {user_id} are')
for recommend in movies.iloc[top_recommends]['title'].values:
    print(recommend)

```

Making Recommendations

The **recommender function** will first display **top 5 highly rated** movies by a given **Customer_ID** and then it'll display **top 5 recommended** movies for this user. **Do not expect these recommended movies to be similar** to these **top 5 highly rated** movies by the user. **Top 5 highly rated** movies are displayed only to give you a gist of this particular user's taste however this meagre list of top 5 cannot capture the his taste completely (Imagine this user has given 5 star rating for 40 movies and the top 5 highly rated movies by this user will be just the first instance of the top 5). But the **top 5 recommendations take into all the movies that the user has rated**.

In [31]: `recommendations(160179)`

The list of top 5 rated movies by the Customer_Id 160179 are
Proof Positive
The Bible Collection: Joseph
Dragon Ball: King Piccolo Saga: Part 2
Last Flight of Noah's Ark
Grind

The list of top 5 recommended movies for the Customer_Id 160179 are
Bonjour Monsieur Shlomi
Lois & Clark: The New Adventures of Superman: Season 1
Kojak: Season 1
It Was a Wonderful Life
Inner Senses

In [32]: `recommendations(82281)`

The list of top 5 rated movies by the Customer_Id 82281 are
Beauty and the Beast: Special Edition
The Velocity of Gary
Love Me Tonight
Head On
Gladiator

The list of top 5 recommended movies for the Customer_Id 82281 are
Para Para Sakura
Disorganized Crime
High Times' Potluck
Worth Winning
Earth vs. The Flying Saucers

Recommending Similar Movies While Searching

```

In [33]: # nearest_neighbor object to find the similar movies
similar_movies=NearestNeighbors(metric='cosine', n_jobs=-1, algorithm='brute')

# fitting the above Learning algo on the utility matrix
similar_movies.fit(utility)

def search():
    # lower casing the searched title
    movie_name=input('Enter your search phrase here: ').lower()

    # finding the matches in the list of movie titles
    filtered=movies[movies.title.str.lower().str.contains(movie_name)]

    if len(filtered)==0:
        print("Couldn't find this movie. Please try again!")
        return
    if len(filtered)>10:
        print("Search phrase is too generic. Please try again!")
        return

    print(f'\nTop matches for your search phrase are:')
    for title in filtered.title.values:
        print(title)

    match_ids=filtered.id.values # list of search phrase matching movie ids
    similar_ids=[]
    seen=set() # to avoid repeatedly recommending same movie

    for ids in match_ids:
        # capturing cosine distance and similar movie id from neighbors
        # for ids from every matching id
        dis,mov=similar_movies.kneighbors(utility[ids],
                                          n_neighbors=6,
                                          return_distance=True)

        for i in range(1,6):
            # adding the cosine distance and similar movie id as long as
            # the similar movie id is not already in the match_ids list
            if (mov[0][i] not in match_ids) and (mov[0][i] not in seen):
                seen.add(mov[0][i])
                similar_ids.append([dis[0][i],mov[0][i]])

    similar_ids.sort(key=lambda x:x[0]) # sorting based on the distance

    print(f'\nRecommended watch:')
    for dist, mov_id in similar_ids[:10]:
        # printing only the titles of the top 10 most similar movies
        print(movies.title.values[movies.id==mov_id][0])

```

In [34]: search()

Enter your search phrase here: sdfhjk
 Couldn't find this movie. Please try again!

In [35]: search()

Enter your search phrase here: a
 Search phrase is too generic. Please try again!

In [36]: search()

Enter your search phrase here: the godfather

Top matches for your search phrase are:

The Godfather Part II
The Godfather Part III
The Godfather
The Godfather Trilogy: Bonus Material

Recommended watch:

GoodFellas: Special Edition
One Flew Over the Cuckoo's Nest
The Silence of the Lambs
Apocalypse Now
Star Wars: Episode I: The Phantom Menace
Hannibal
Scarface: 20th Anniversary Edition
The Devil's Advocate
Seven: Bonus Material
Scarface: 20th Anniversary Edition: Bonus Material

In [37]: search()

Enter your search phrase here: finding nemo

Top matches for your search phrase are:

Finding Nemo (Widescreen)
Finding Nemo (Full-screen)

Recommended watch:

Monsters Inc.
Shrek (Full-screen)
Shrek 2
A Bug's Life
The Incredibles
Toy Story 2

In [38]: search()

Enter your search phrase here: Shaolin Soccer

Top matches for your search phrase are:

Shaolin Soccer

Recommended watch:

Kung Fu Hustle
The Legend of Drunken Master
Iron Monkey
The Blind Swordsman: Zatoichi
Once Upon a Time in China 2

You can see from the above examples that

If we search for a **gangster movie, then our model **recommends** similar **gangster movies**.**

If we search for a **animated movie, then our model **recommends** similar **animated movies**.**

If we search for a **chinese movie, then our model **recommends** similar **chinese movies**.**

The highlight is that, all these **recommendations** are purely based on the **similarity** between the **ratings** of the movies, not based on the **similarity** between the **content** of the movies. Because our **model** is built using **collaborative filtering** approach and not using **content based filtering** approach.

Hope you found this notebook useful!

Lets connect

<https://www.linkedin.com/in/fazil-mohammed-4062711b2/> (<https://www.linkedin.com/in/fazil-mohammed-4062711b2/>)