



What's Slowing You Down?

Using Kernel Event Tracing to Uncover Performance Issues

Malte Mundt, Senior Field Application Engineer, BlackBerry QNX
mmundt@blackberry.com

Agenda

01

Overview: What does Kernel
Event Tracing mean?

02

System Profiler: Visualization
and Interpretation

03

Finding out about...
Processing Time

04

Finding out about...
Dependencies

05

Finding out about...
Bottlenecks

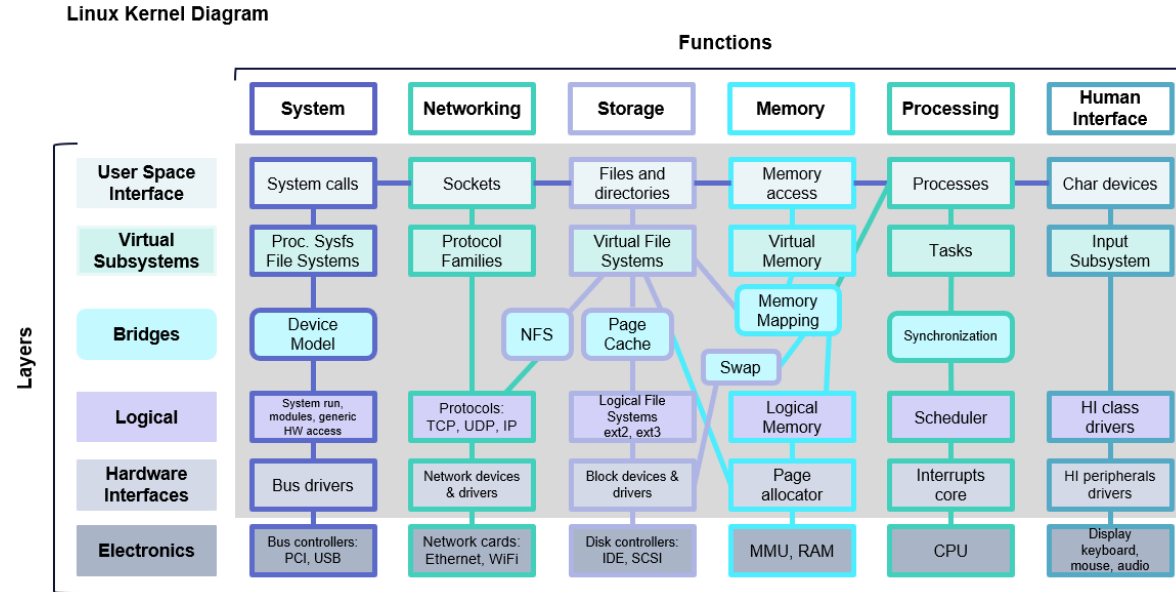
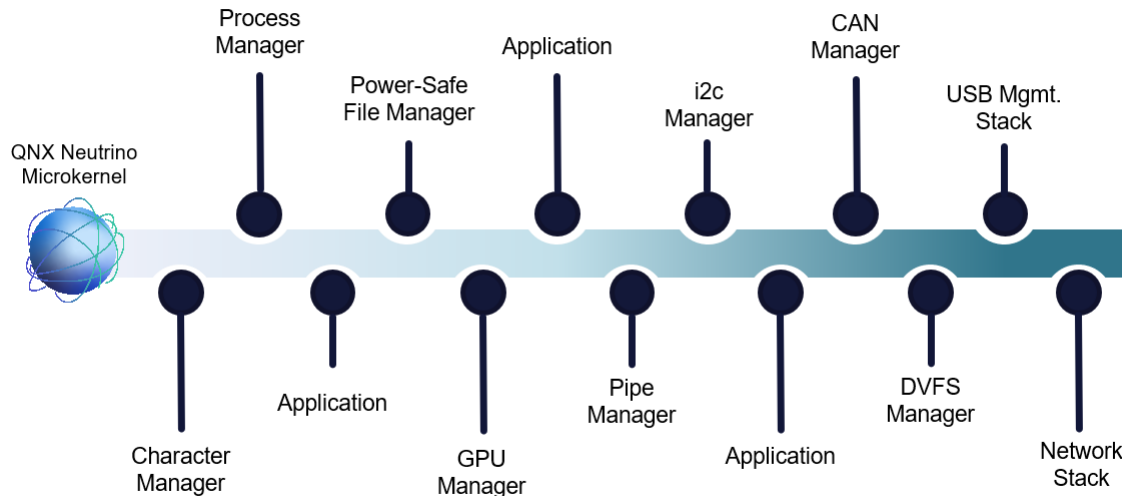
06

Conclusion
Q&A

What does Kernel Event Tracing mean?

Overview

What Kernel Tracing means in a Microkernel environment



- With the QNX Microkernel, everything is a process with threads: Drivers, Stacks, Services, Applications
- All processes can use the same APIs, i.e. ANSI C / C99 / C11, POSIX PSE54 etc.
- Kernel Event Tracing is logging kernel events, System Profiling visualizes kernel event log
- The QNX microkernel can be extended to become a Hypervisor: VMs are special processes

Performance analysis from the kernel's point of view

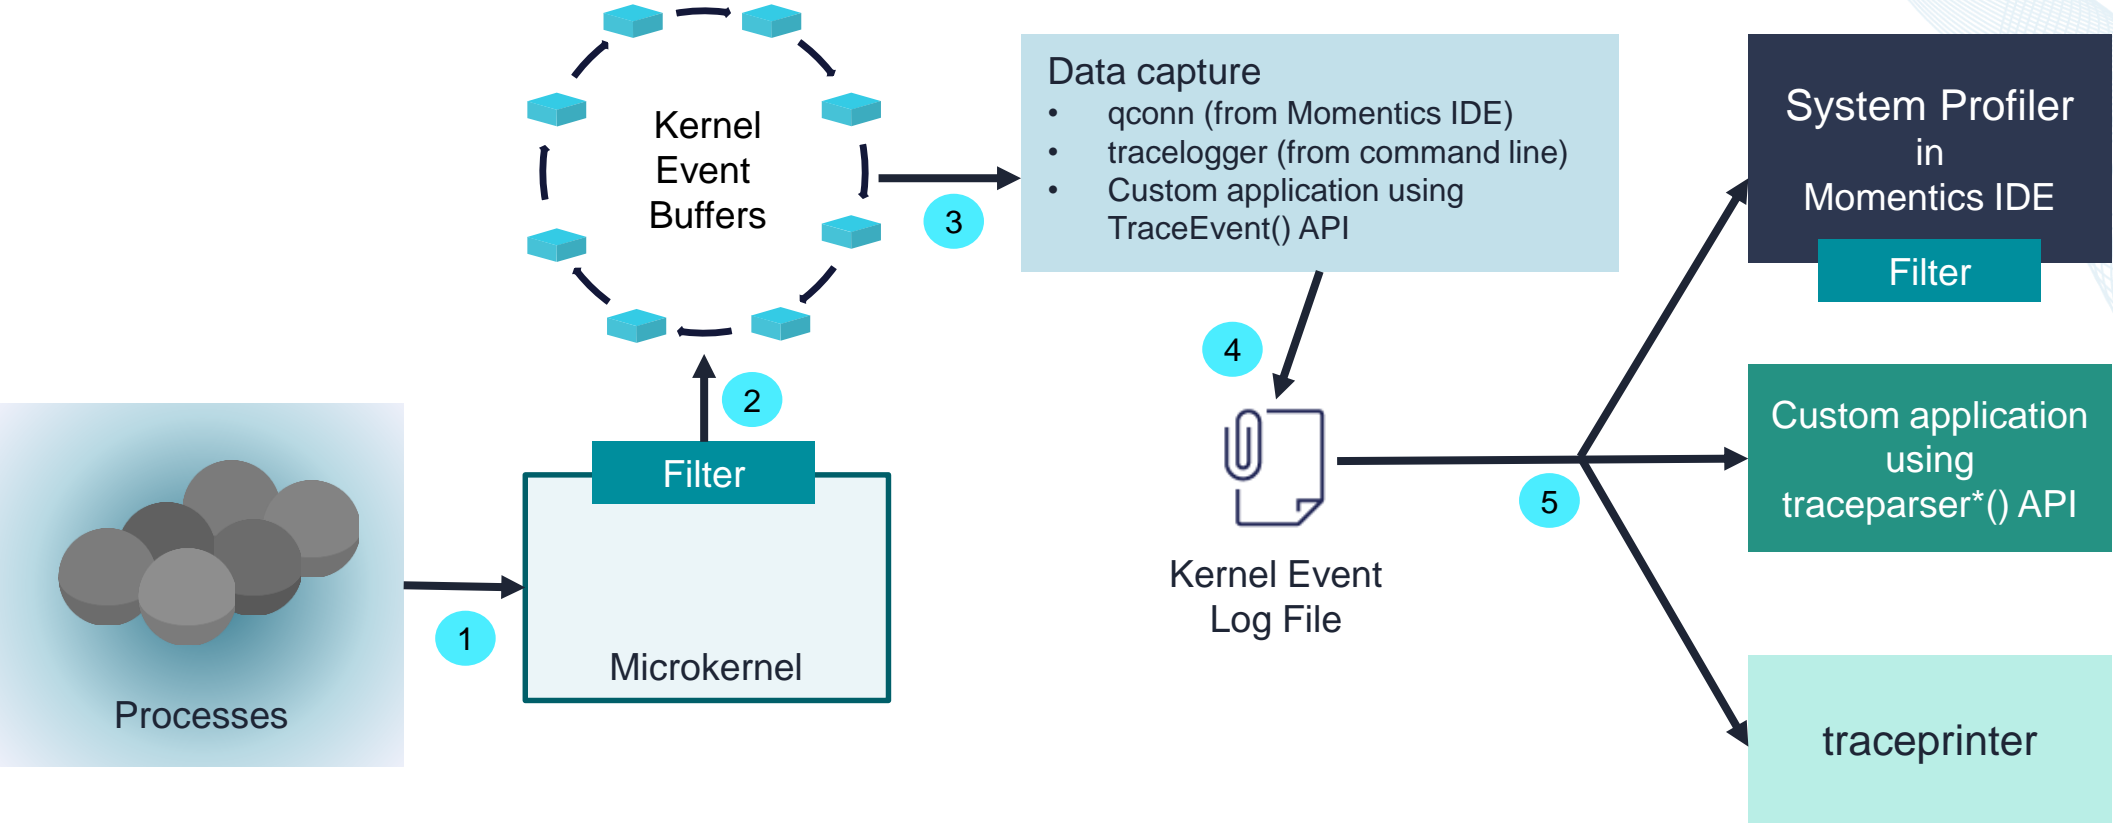
Kernel Event Tracing... is using the instrumentation built into the QNX kernel for logging of events:

- Creation and Exit of Threads, Processes, Virtual Machines
- Kernel calls made by Threads
- Message Passing based IPC, Thread Synchronization
- Thread State Changes and Scheduling
- Guest Exit/Guest Entry (Hypervisor)
- Interrupt Handling
- User-defined trace events

... gives you a holistic view on overall system performance

... lets you drill down to the details

QNX Kernel Event Tracing Data Flow

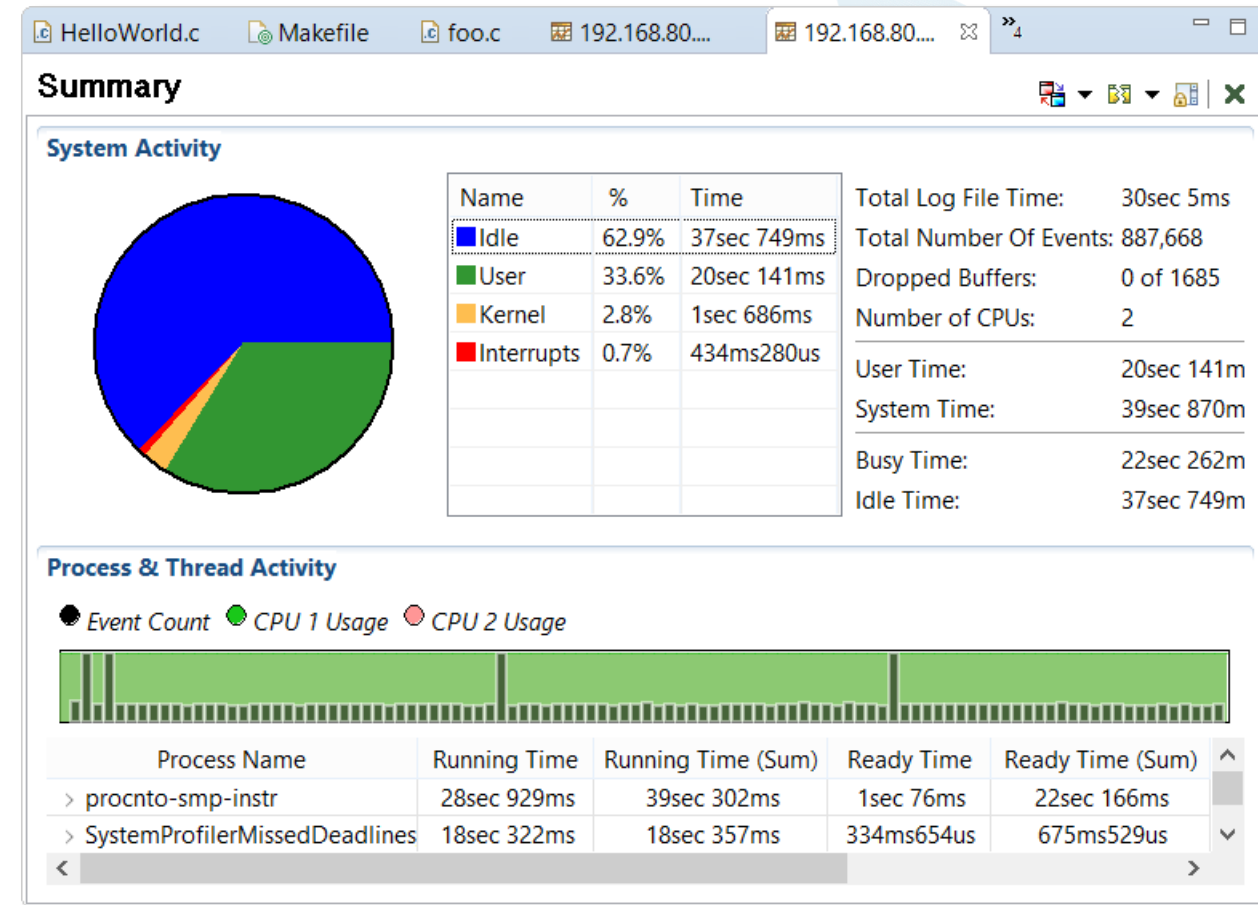


Kernel Event Tracing - Visualization and Interpretation



System Profiler

Find out overall CPU usage level

- Events have microsecond resolution timestamps
 - Number of events displayed in bar graph
- System Profiler calculates CPU time
 - Idle
 - User Space processes
 - Kernel
 - Interrupt handlers
- Most active processes also shown in Summary



Identify sources of CPU load

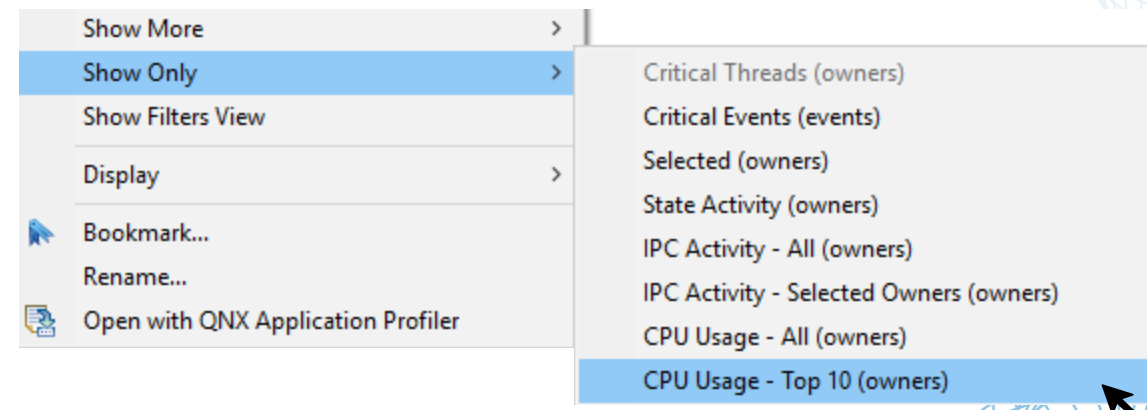
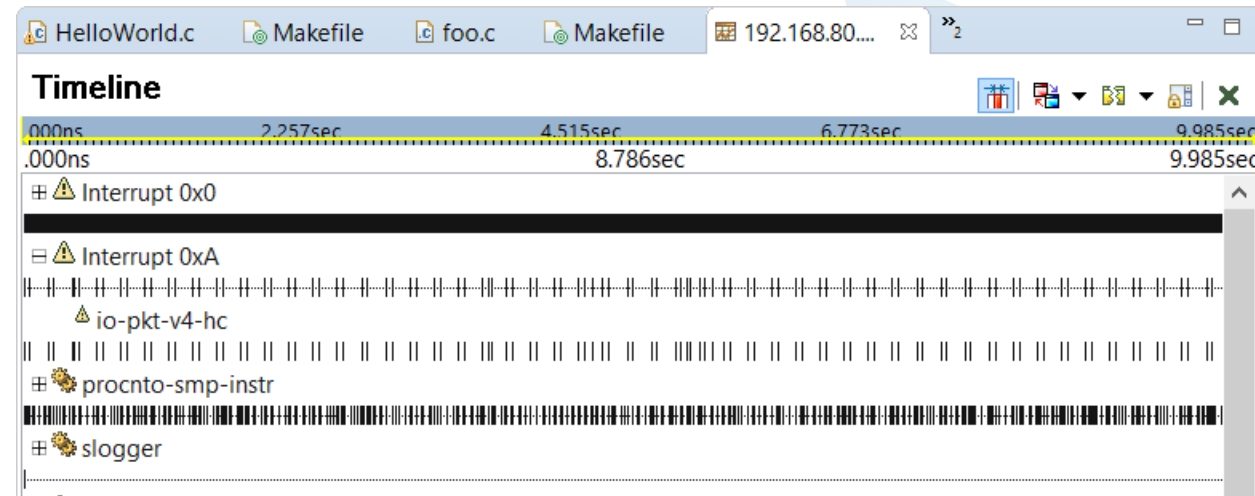


- Microsecond-accurate CPU activity graph per core – identify peaks, verify actual usage
- Zoom in to see CPU toggling between active and idle

- Identify processes causing base load, peaks
- Select/deselect individual threads to see CPU load accumulating

Maximum Detail: The Time Microscope

- Timeline view sorts all events by source (processes and threads) to represent their exact timing
- For each event source a timeline is drawn
 - Events represented by vertical tick marks
 - Interrupts listed at the top
- Large systems can have hundreds of processes, thousand of threads – filtering is key
 - Filter out unneeded event classes when logging
 - Event Display filtering: Show only top 10 CPU users, filter certain process names, event classes...

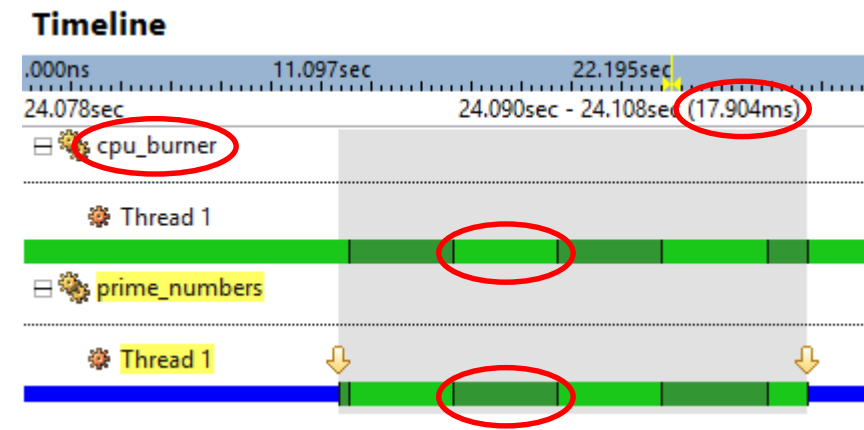
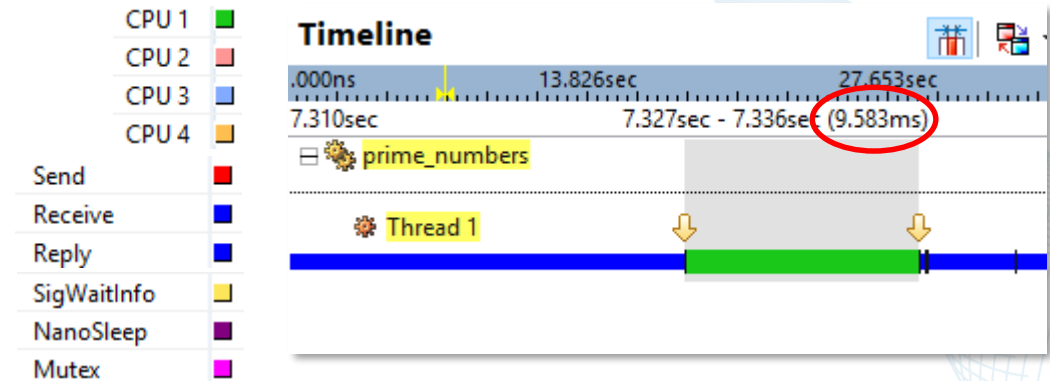


What can you find out with a Kernel Event Trace?

Processing Time




Exact measurement of processing times

- Threads that want to use CPU are *Ready*
 - Scheduler makes them *Running* on available cores
- Threads waiting for something are *Blocked*
 - Color shows why the thread is not running
 - All thread states are documented
- Example: An event happens, you have to quickly calculate something – Timeline shows you exactly how long it takes
- But when CPU is fully loaded...
 - ... suddenly your calculation takes much longer
- Timeline over multiple threads helps understanding why



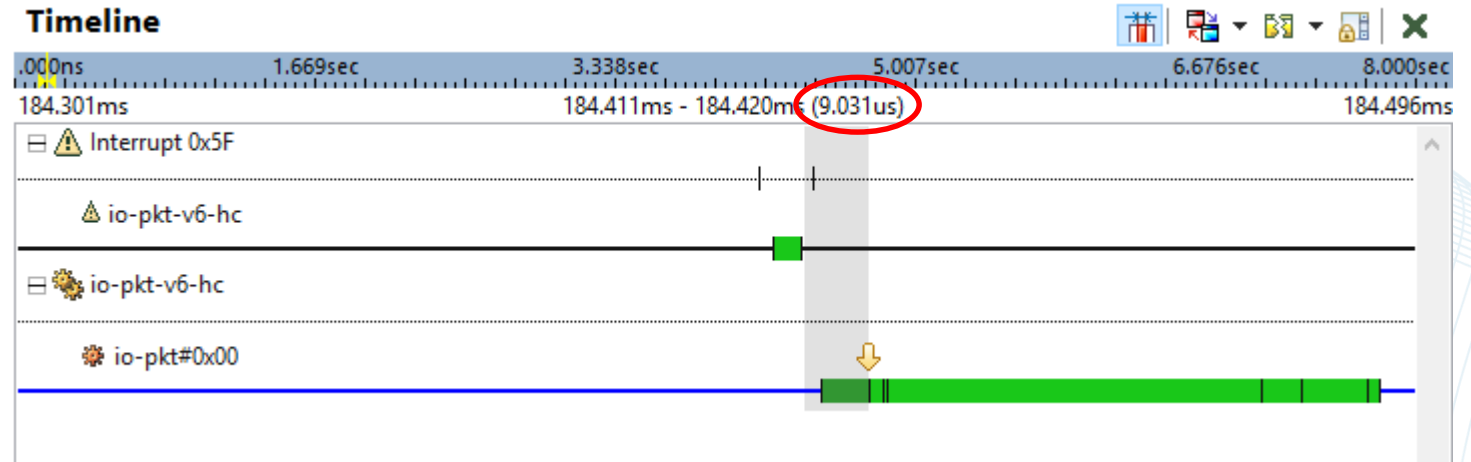
Analyze scheduler actions, adjust priorities

- More threads *Ready* – calculation slower...
- With Round Robin only, Ready Queue can get long
- Scheduler evaluates thread Priority first
- Let's increase the priority of our calculation
 - Now it's back to its original speed
- Same total load, drastically different behavior



Interrupt frequency and handling times

- Find out actual frequency of interrupts
- How long until handler gets called?
- IRQ handler execution time?
- Thread to be scheduled?




- How long does it take until the thread is made *Running* ? Depends on priority...
- Thread woken up on behalf of interrupt handler is executing in full POSIX environment
 - On any available core and in your configured priority
- Optimize? Use only a handler or only a thread

What can you find out with a Kernel Event Trace?

Dependencies



Performance dependencies – Clients and Servers

- Our CPU load example prints to console when done – it's IPC
- Visualize Message Passing based IPC to see dependencies
 - Servers and clients are both processes, interaction is IPC
 - Drivers and Stacks (Network, USB, Audio) are server processes
- Application performance depends on servers they are interacting with
 - CPU load of servers induced by their clients
- System Profiler enables you to identify and measure impact of servers on your application
 - Measure how long your thread is blocked when it's a client waiting for a server
 - Look at server threads' processing time
 - Find other CPU-heavy threads that slow down server



Visualize and analyze IPC interactions


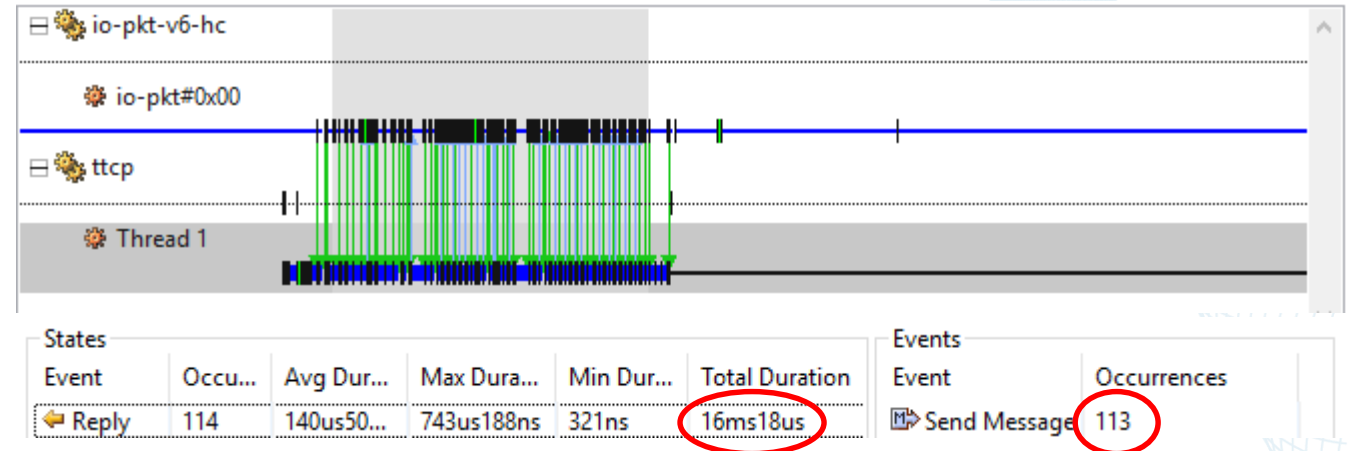
- Example: Application sends data out via TCP/IP
- Select process of interest, filter on IPC partners
- Our test app is client of various servers
 - Heavy dependency on network stack is obvious
- Server spends a lot of time for application
- Client spends a lot of time waiting for sever
- Client/Server CPU statistics reveal *Imposed Time*



Client/Server CPU Statistics			
Data of tracebuffer.kev			
Owner	Total Time	Self Time	Imposed Time
ttcp - Thread 1	23us432ns	2us272ns	21us159ns
io-pkt-v6-hc - io-pkt#0x00	24us402ns	24us402ns	0ns

Identify IPC frequency, Client wait times

- Client needs time to prepare buffer
- Server process time usually is sum of
 - Fixed time needed to process request
 - Variable time depending on the size of the data
 - Both include hardware interaction
- Statistics View shows Client wait time
- Determine Optimization options:
 - Server processing, Client calls, IPC parameters, Shared Memory




What can you find out with a Kernel Event Trace?

Bottlenecks

Hey, we are waiting for you! The absence of full CPU load

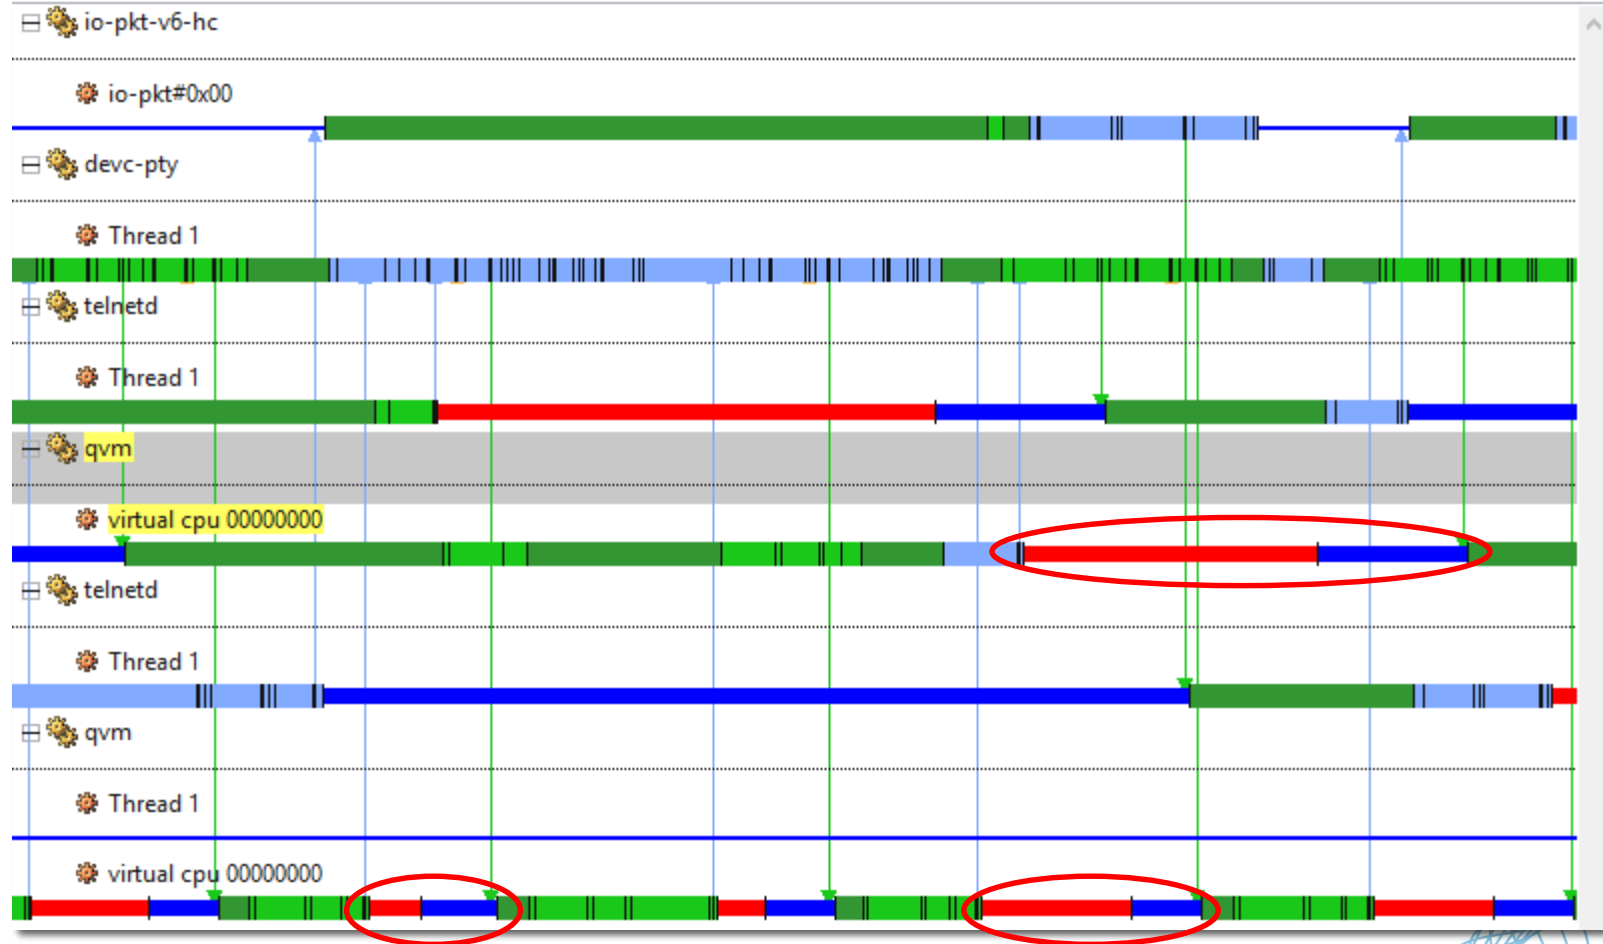
- Scenario: Something is too slow (i.e. button press, reaction takes 2s) – but CPU is not fully loaded
- Resource contention may slow you down – waiting for a Mutex, Condvar, Semaphore
- System Profiler will help you identify:
 - Wait time – Average and Maximum duration
 - Threads which are holding Mutex longest while others wait
 - Unveil thread that does not release Mutex at all (Deadlock!)

States					
Event	Occu...	Avg Duration	Max Duration	Min Duration	Total Duration
✱ Running	5	85us969ns	104us575ns	37us651ns	429us863ns
🚦 Ready	5	15us121ns	18us90ns	13us833ns	75us606ns
🔒 Mutex	5	292us515ns	1ms206us	51us560ns	1ms462us
⌚ NanoSleep	1	21us757ns	21us757ns	21us757ns	21us757ns



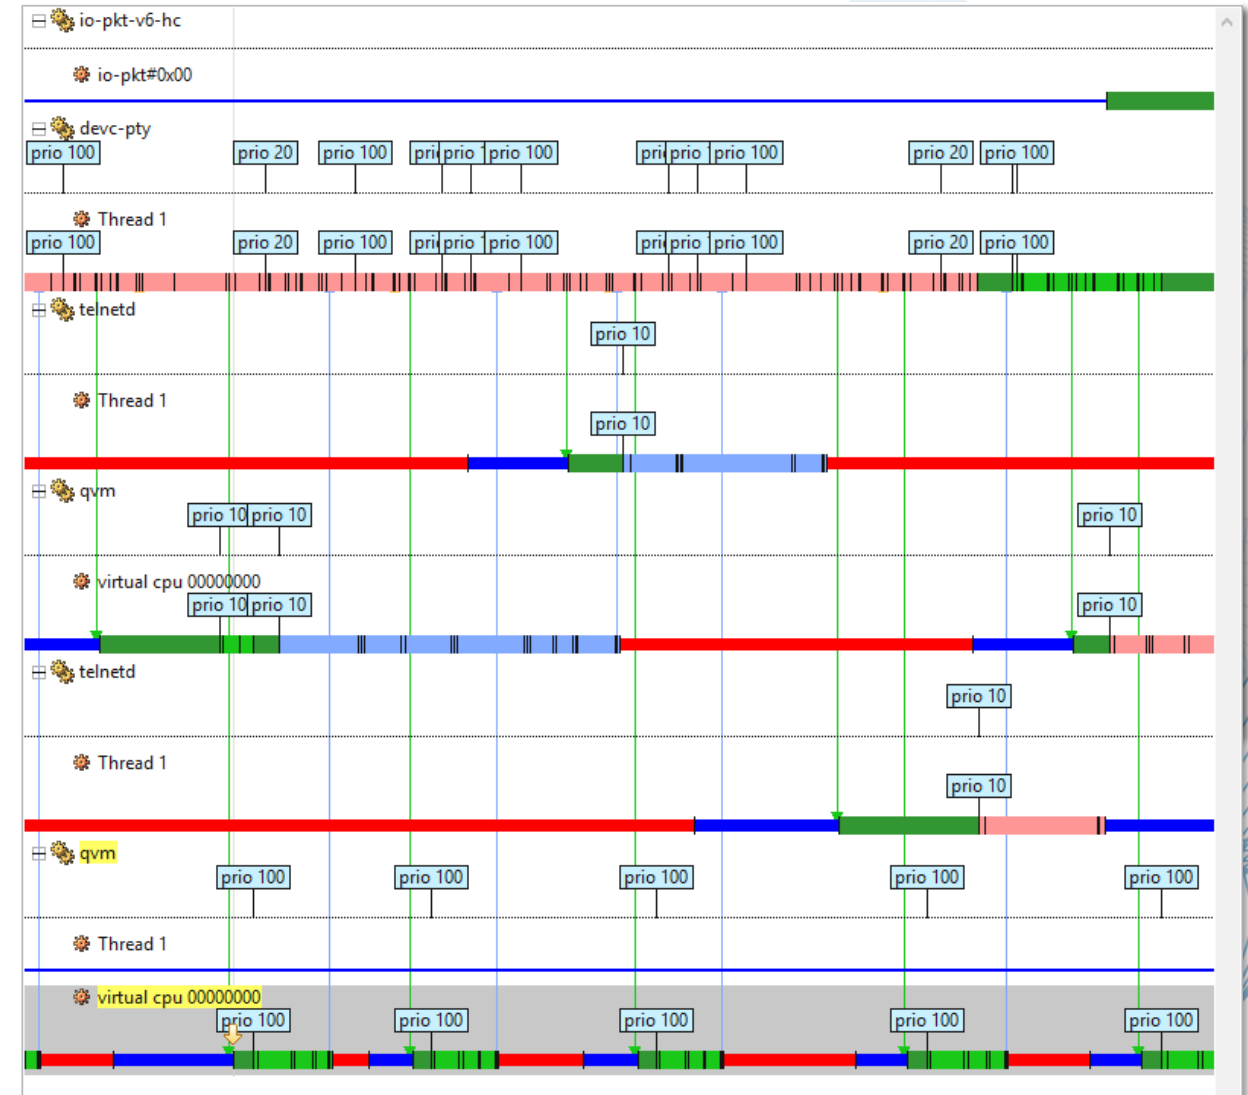
Servers can be bottlenecks – how priority inheritance helps

- Example: QNX Hypervisor
- I telnet in twice, and from each console I start a virtual machine
 - Linux guest at priority 10
 - QNX guest at priority 100
- Linux guest prints out text to console ...
- ... then QNX guest also prints text simultaneously
- Servers managing shared resources will inherit client priority – wait time minimized



Unveiling another kind of bottleneck

- Everyone seems busy – but CPU not fully loaded
- Long red lines are eye-catching: Send-Blocked!
 - Meaning: Receiver busy, can not receive
- Multiple threads are Send-Blocked
- Priority Labels show scheduling priority...
 - “devc-pty” priority is changing while it’s running
 - It is bombarded with requests
 - It is running the entire time
 - It’s single-threaded! 🤪



What's Slowing You Down?

Conclusion

What's Slowing You Down? – Conclusion

Processing Time

- Sources of CPU load
- Priorities, Ready Queue
- Interrupts

Dependencies

- Client/Server Relationships
- Imposed Time
- IPC Frequency

Bottlenecks

- Threads waiting for each other
- Priority Inheritance
- Send-Blocked clients

Kernel Event Tracing and System Profiling support you to:

- Understand your system, realize how all its components are intertwined and play together
- Inspect any individual processes regarding its CPU usage, IPC, priorities, resource contentions
- Make the right decisions re performance improvement and optimization strategies

Q&A

Malte Mundt
Senior Field Application Engineer
mmundt@blackberry.com

Thank you!